

Assignment 3: Agentic RAG System with Azure OpenAI & LangGraph

1. Background

An **Agentic Retrieval-Augmented Generation (RAG)** system combines:

- A **Knowledge Base (KB)** indexed in a vector database,
- An **LLM** (for generating answers),
- A **self-critique loop** (to check completeness),
- A **refinement step** (to fill missing gaps).

This setup ensures the system can answer **general questions** and also **leverage what's in the vector database** to provide citation-backed responses.

2. Problem Statement

“Build an Agentic RAG system on Azure using LangGraph that retrieves up to 5 KB snippets, generates an answer with Azure GPT-4 mini, critiques it, and when required, refines it with one additional snippet from Pinecone or Weaviate (trial cloud instance). The system should log outputs and support observability with MLflow.”

Dataset: self_critique_loop_dataset.json

Dataset and a pinecone starter notebook is also available on LMS in Assignment 3 files.

3. Detailed Tasks

1. Preprocessing & Indexing

- Load KB JSON (~30 entries).
- Use **Azure Embeddings**: text-embedding-3-small.
- Store vectors in **Pinecone** or **Weaviate trial cloud instance** (no local hosting).

2. LangGraph Workflow Define 4 nodes:

- **Retriever Node**: fetch top-5 snippets.
- **LLM Answer Node**: use **Azure GPT-4 mini** (temperature=0) to generate initial answer with [KBxxx] citations.
- **Self-Critique Node**: Gemini checks if answer is COMPLETE or needs REFINE.

- **Refinement Node:** if REFINE, retrieve 1 more snippet and regenerate answer.

Decision logic:

- If COMPLETE → return initial answer.
- If REFINE → return refined answer.

3. Tracing & Observability

- Integrate **MLflow** to log runs, retrieved snippets, model outputs, critique results, and final answers.
-

4. Tools & Tech

- **Azure GPT-4 mini** (deployed in Azure OpenAI)
- **Azure Embeddings** (text-embedding-3-small)
- **Vector DB:** Pinecone (trial) or Weaviate (trial)
- **LangGraph** for pipeline wiring
- **MLflow** for observability
- **Python 3.10**
- Suggested packages:

```
langgraph
azure-ai-inference
pinecone-client    # or weaviate-client
mlflow
pydantic
```

5. Sample Queries

1. “What are best practices for caching?”
 2. “How should I set up CI/CD pipelines?”
 3. “What are performance tuning tips?”
 4. “How do I version my APIs?”
 5. “What should I consider for error handling?”
-

6. Deliverables

Submit either:

- **Jupyter Notebook** with all steps, OR
 - **ZIP folder** with:
 - `index_kb.py` (embeddings + vector DB indexing)
 - `agentic_rag_azure.py` (LangGraph workflow + MLflow logging)
 - `requirements.txt`
-

7. Notes

- Always cite snippets [KBxxx].
 - Keep flow simple: 1 critique, max 1 refinement.
 - Use trial vector DB instances (Pinecone/Weaviate portal).
 - Temperature=0 for deterministic outputs.
-